

LLM-Powered Query Monitoring and Optimization for Reproducible Big Data Workloads

Shouvik Sharma

K. J. Somaiya College of Engineering

India

Email: shouvik.s@somaiya.edu

Abstract—Big data platforms execute heterogeneous analytical SQL workloads across scientific, operational, and security datasets. Inefficient queries can increase cost, delay downstream analytics, and reduce trust in shared data infrastructure. This paper presents a reproducible framework for LLM-powered query monitoring and optimization. The framework ingests publicly accessible datasets, executes a controlled DuckDB workload, stores query telemetry in SQLite, analyzes SQL with a large language model, records optimization recommendations, and validates rewritten queries through execution-based result checks. In an evaluation over five logical datasets and 91 queries (41 baseline, 50 inefficient) targeting 12 distinct anti-pattern types, the framework achieves 96.9% detection accuracy, 0% false positive rate, 93.8% recall, and a 93.3% tested-instance result-equivalence rate for rewrites, with total LLM API cost of \$0.005522. Compared against *squaudit*, a rule-based static analyzer with 27 built-in rules, the LLM-based approach eliminates false positives through contextual understanding of query intent while matching detection recall. The framework and its evaluation artifact are publicly available, providing a foundation for reproducible research on foundation-model-assisted query governance in big data systems.

Index Terms—big data systems, query monitoring, SQL optimization, large language models, reproducibility, benchmark artifacts

I. INTRODUCTION

Big data systems are increasingly used to integrate scientific, environmental, cybersecurity, and operational datasets. As these platforms grow, SQL workloads become more diverse: analysts issue ad hoc queries, dashboards run recurring aggregations, and data applications generate programmatic SQL. Inefficient queries can silently increase runtime, waste compute, and produce avoidable operational cost. Query monitoring is therefore a governance problem as well as a performance problem.

Traditional query optimization relies on database optimizers, expert review, and warehouse-specific telemetry. These mechanisms remain essential, but they do not fully address the human-facing problem of explaining why a query is risky, identifying common anti-patterns, and proposing understandable rewrites. Recent foundation models create an opportunity to augment query review with natural-language reasoning and code transformation. This paper addresses this challenge by embedding LLM analysis within an execution-based validation harness that confirms rewrite correctness through re-execution and result comparison.

This paper presents an LLM-powered query monitoring framework designed for reproducible experimentation. The system downloads public datasets, constructs a controlled workload of 91 queries spanning five datasets and 12 anti-pattern types, executes queries against DuckDB [1], records telemetry in a SQLite query-history database [2], analyzes SQL with an LLM, stores recommendations, executes rewritten queries, and reports detection, result-equivalence, runtime, and cost metrics. The framework is open source [3].

The paper contributes: (1) a modular query-monitoring architecture for foundation-model-assisted SQL review; (2) a reproducible artifact linking public datasets, workload execution, query history, LLM analysis, and reporting; (3) a rigorous evaluation methodology with statistical analysis of detection accuracy, result equivalence, and cost; and (4) an empirical comparison against rule-based static analysis demonstrating the advantage of contextual understanding.

This work addresses four research questions. **RQ1:** How accurately does LLM-assisted review detect SQL anti-patterns on a labeled multi-domain workload? **RQ2:** How often do generated rewrites execute successfully and preserve tested-instance results? **RQ3:** What latency and monetary-cost overheads are introduced by LLM-assisted monitoring? **RQ4:** How does performance vary across datasets and anti-pattern categories?

II. RELATED WORK

We organize prior work into four themes that jointly motivate the proposed framework and highlight the gap this work addresses.

A. SQL Anti-pattern Detection and Static Analysis

A mature ecosystem of static analysis tools exists for identifying risky SQL patterns through deterministic rule-based checks. Tools such as *squaudit* [4] detect SELECT *, cross joins, missing WHERE clauses, implicit type coercions, and invalid predicate pushdowns from raw query text. At warehouse scale, Google BigQuery Anti-pattern Recognition [5] scans INFORMATION_SCHEMA for top-slot-consuming jobs, flagging missing partition filters and suboptimal join order through execution statistics. DWH-Auditor [6] similarly detects full table scans, zombie tables, and under-utilized materialized views in enterprise data warehouses. These tools share three limitations: (1) they operate

on fixed rule sets and cannot adapt to context-dependent inefficiencies that require semantic understanding of the data domain; (2) they flag anti-patterns without proposing concrete rewrites, leaving remediation to the user; and (3) they provide no mechanism to validate whether a proposed fix preserves query results on the tested instance. Our work addresses these gaps through LLM-based semantic analysis for detection and rewrite generation, coupled with execution-based tested-instance validation for result checking.

B. LLMs for SQL Optimization and Rewriting

Recent work demonstrates that LLMs can understand SQL semantics and suggest meaningful optimizations. R-Bot [7] proposes a multi-source rewrite evidence pipeline with hybrid structure-semantics retrieval, achieving state-of-the-art rewrite quality on production Huawei workloads. LLMOpt [8] fine-tunes LLMs for plan candidate generation, outperforming traditional optimizers on the Join Order Benchmark by exploring a broader plan space. LLMSTEER [9] shows that LLM embeddings of raw SQL contain sufficient semantic information for a binary classifier to outperform heuristics on plan selection, suggesting that even frozen LLM representations encode query properties relevant to optimization. LaPuda [10] introduces a policy-based multi-modal optimizer with a cost descent algorithm that learns from execution feedback. LLM4Hint [11] demonstrates that moderate-sized LLMs (7B parameters) can recommend effective optimizer hints, reducing the need for large models. Despite these advances, existing LLM-based approaches focus narrowly on optimizer hint selection or plan space exploration. No existing system jointly persists query telemetry, records structured LLM recommendations, validates generated rewrites through re-execution, and reports per-query model cost.

C. Cloud Warehouse Monitoring and Governance

Production cloud data warehouses have motivated significant work on workload monitoring and resource governance. SnowPatrol [12] applies unsupervised anomaly detection to Snowflake usage metadata, identifying outlier queries by runtime and slot consumption. Keebo [13] continuously monitors warehouse performance telemetry to automate sizing decisions and auto-suspend intervals, reducing idle compute by up to 40%. Bress et al. [14] conducted a large-scale empirical study of 667 million production queries from Snowflake, identifying that 62% of queries access fewer than 10 columns and that missing filters are among the most common performance issues. Patterns [15] extracts historical query patterns to recommend partitioning keys and clustering orders via AI-driven analysis of query access paths. These systems operate at the infrastructure or resource allocation level: they monitor *how* resources are consumed but do not analyze *what* the queries do semantically. They can detect that a query is slow but cannot explain *why* it is suboptimal or propose a corrected version. Our framework bridges this gap by combining lightweight telemetry capture (runtime, row counts, checksums) with

LLM-driven semantic analysis that produces both explanations and executable rewrites.

D. Benchmarks for Reproducible Evaluation

The database community has established several benchmarks for evaluating query processing and optimization. TPC-H [16] and TPC-DS [17] remain the gold standard for measuring analytical query throughput and resource utilization across standardized schemas and data distributions. The Join Order Benchmark (JOB) [18] provides real-world queries derived from the IMDB dataset, capturing the complexity of multi-way joins with realistic cardinality estimation challenges. BIRD-INTERACT [19] introduces a dynamic interaction framework for text-to-SQL evaluation, where the model can issue database exploration queries before generating a final answer. SPIDER [20] and its variants provide cross-domain text-to-SQL benchmarks with annotated queries. However, none of these benchmarks targets the specific task of evaluating LLM-based query monitoring frameworks. A suitable benchmark requires: (a) labeled workloads containing both optimal and intentionally inefficient query variants with ground-truth annotations; (b) multiple datasets spanning diverse domains to test cross-domain generalization; (c) execution telemetry (runtime, cardinality, cost) for both original and rewritten queries; and (d) standardized metrics for detection accuracy and result-equivalence rate. This paper contributes a labeled workload and evaluation methodology that addresses these requirements.

E. Gap Statement

While prior work has explored LLM-based query optimization [7]–[11] and static SQL analysis [4]–[6], our framework combines four elements in a single auditable artifact:

- 1) **Persistence-first governance:** All query telemetry, LLM recommendations, and validation results are stored in a structured QueryHistoryStore, enabling auditable provenance and replayability.
- 2) **Hybrid detection:** Deterministic static analysis flags known anti-patterns, while LLM-based review identifies contextual risks (e.g., missing join keys in complex queries).
- 3) **Execution-based validation:** Generated rewrites are validated through re-execution and result comparison, checking tested-instance result-equivalence.
- 4) **Cost-aware monitoring:** Per-query LLM API costs are tracked and reported, enabling tradeoff analysis between accuracy and expense.

This end-to-end pipeline addresses gaps in prior systems, which either lack persistence, rely solely on static rules, or omit execution validation.

III. SYSTEM DESIGN AND METHODOLOGY

The framework implements a modular pipeline that transforms raw SQL queries into structured, validated, and cost-accounted optimization recommendations. Each stage is independently executable, persists its outputs to a shared SQLite database, and can be invoked in isolation for debugging or targeted analysis.

TABLE I
RULE-BASED VS. LLM-ASSISTED COVERAGE

Anti-pattern	Rule-based detection	LLM-assisted advantage
SELECT *	Yes, via regex or AST rules	Explains why projection breadth matters for the workload
Missing predicate	Yes, via query-structure checks	Can distinguish benign filters from risky full scans
Cross join	Yes, via join-graph analysis	Can suggest a safe rewrite and describe the cardinality impact
Redundant sub-query	Limited	Can detect unnecessary nesting and recommend flattening




Fig. 1. Six-stage pipeline architecture. Arrows indicate the primary data flow direction; each stage reads from and writes to the shared QueryHistoryStore.

A. System Overview and Data Flow

Figure 1 illustrates the six-stage pipeline through which a query travels. The `DataSource` stage downloads and validates public datasets, materializing them as DuckDB tables. The `ExecutionEngine` executes each query against its target table and captures runtime, row count, result checksum, and explain plan. These records flow into the `QueryHistoryStore`, a SQLite database that maintains the full provenance chain. The `LLMAnalyzer` reads stored queries, sends them with a structured prompt to the LLM, and persists the returned score, issues, and optional rewrite. For flagged queries (score $\geq \tau$), the `RecommendationEngine` executes the rewritten query through the same harness, validates tested-instance result equivalence, and computes cost deltas. Finally, the `ReportGenerator` aggregates all results into detection accuracy, result-equivalence rate, and cost summaries. Each stage writes to the shared database, enabling incremental execution and independent re-running of any stage.

Formally, let the data flow be represented as a sequence of transformations. Starting with a dataset $d \in D$ and workload $Q = \{q_1, \dots, q_n\}$:

$$\text{Stage 1 (DataSource)} : d \mapsto \text{table}(d) \quad (1)$$

$$\text{Stage 2 (Execution)} : (q_i, \text{table}(d), e) \mapsto M_i \quad (2)$$

$$\text{Stage 3 (Storage)} : (q_i, M_i) \mapsto \text{DB.executions} \quad (3)$$

$$\text{Stage 4 (LLM Analysis)} : q_i \mapsto (s_i, I_i, q'_i) \quad (4)$$

$$\text{Stage 5 (Recommendation)} : (q'_i, s_i \geq \tau) \mapsto M'_i, \text{resEq}_i \quad (5)$$

$$\text{Stage 6 (Reporting)} : \{(M_i, M'_i, \text{resEq}_i)\} \mapsto \text{metrics} \quad (6)$$

The pipeline is designed to be replayable and auditable. Deterministic stages such as dataset staging, query execution, validation, and report generation can be re-run with the same inputs. External LLM calls are not inherently idempotent

TABLE II
TARGET ANTI-PATTERN TAXONOMY

Anti-pattern	Formal Characterization	Expected Impact
SELECT *	Projects all columns when only a subset is needed	More I/O, wider shuffle, higher memory use
Missing predicate	Scans all rows without a filter condition	Full scan, excess I/O
Cross join	Joins tables without a join predicate	Cartesian product, many intermediate rows
Redundant subquery	Wraps a simple query in unnecessary nesting	Blocks predicate pushdown

because hosted-model responses can vary across repeated calls, backend changes, or provider revisions. The framework therefore persists prompts, model identifiers, scores, token counts, recommendations, rewrites, and validation outputs so prior analyses can be replayed from stored records.

B. Problem Formulation

Let a **query workload** be a set of SQL queries $Q = \{q_1, \dots, q_n\}$ executed against a dataset $d \in D$ on an engine $e \in E$. The **execution function**

$$\text{exec}(q_i, d, e) \mapsto M_i = (t_i, r_i, c_i, p_i) \quad (7)$$

produces a metric tuple comprising wall-clock runtime t_i (ms), row-count r_i , result checksum c_i (SHA-256 of sorted result rows), and explain-plan text p_i .

The **LLM analysis function**

$$f_{\text{LLM}}(q_i) \mapsto (s_i, I_i, q'_i) \quad (8)$$

assigns a risk score $s_i \in [0, 100]$, identifies issues $I_i = \{i_1, \dots, i_k\}$, and optionally produces a rewritten query q'_i when $s_i \geq \tau$, where $\tau = 40$ is the detection threshold.

The **result validation** function compares original and rewritten executions:

We define $\text{resEq}(q_i, q'_i)$ as 1 when both the row count and sorted result checksum of the rewritten query match the original execution, and 0 otherwise. A rewrite is considered equivalent on the tested instance only when both conditions hold.

The **total LLM API cost** is

$$C_{\text{LLM}} = \sum_{i=1}^n (p_{\text{in}} \cdot \text{tok}_{\text{in}}^{(i)} + p_{\text{out}} \cdot \text{tok}_{\text{out}}^{(i)}) \quad (9)$$

where $p_{\text{in}} = \$0.15/10^6$ tokens and $p_{\text{out}} = \$0.60/10^6$ tokens (GPT-4o-mini) [21]. We report runtime and cost as separate quantities. Standard detection metrics are: TP = inefficient queries scored $\geq \tau$; TN = baselines scored $< \tau$; FP = baselines incorrectly flagged; FN = inefficient queries missed. The framework's taxonomy begins with the four most prevalent SQL anti-patterns reported in production big data workloads [5], [14], and extends to 12 distinct anti-pattern types in the experimental workload (Section V-B). Table II summarizes the core four used to define the formalism.

Algorithm 1 LLM-Powered Query Monitoring Pipeline

Input: Dataset list D , workload Q , engine E , threshold τ
 $R \leftarrow \emptyset$, $C \leftarrow \emptyset$
for each $(q_i, d_i) \in (Q, D)$ **do**
 $M_i \leftarrow \text{exec}(q_i, d_i, E)$ ▷ Execute original
 $\text{store_query}(q_i)$
 $\text{store_execution}(q_i, M_i, \text{label} = \text{original})$
 $(s_i, I_i, q'_i) \leftarrow \text{LLM_analyze}(q_i)$
 $\text{store_llm_analysis}(q_i, s_i, I_i)$
 if $s_i \geq \tau$ **then**
 $M'_i \leftarrow \text{exec}(q'_i, d_i, E)$
 $\text{store_execution}(q'_i, M'_i, \text{label} = \text{rewritten})$
 $\text{resEq}_i \leftarrow \text{resEq}(q_i, q'_i)$
 $\Delta_i \leftarrow \text{cost_delta}(M_i, M'_i)$
 $c_i \leftarrow \text{llm_cost}(\text{tok}_{\text{in}}, \text{tok}_{\text{out}})$
 $\text{store_recommendation}(q_i, q'_i)$
 $\text{store_comparison}(q_i, \text{resEq}_i, \Delta_i, c_i)$
 end if
 $R \leftarrow R \cup \{(q_i, s_i \geq \tau)\}$
end for
return R, C —detection results and cost records are persisted to QueryHistoryStore.

TABLE III
ARCHITECTURAL COMPONENTS AND THEIR DATABASE INTERACTIONS

Component	Responsibility	Database Tables
Data source	Downloads and stages datasets.	datasets
Execution engine	Runs SQL and records runtime and checksums.	queries; query executions
History store	Maintains the SQLite system of record.	all tables
LLM analyzer	Scores queries and stores model output.	queries; LLM analyses
Recommender	Rewrites, validates, and measures changes.	recommendations; cost comparisons
Report generator	Summarizes accuracy and cost.	read-only

Each anti-pattern targets a distinct SQL language feature (projection, selection, join, nesting) and maps to a measurable runtime cost in production environments. The framework’s anti-pattern set is extensible: adding a new pattern requires only adding corresponding query variants to the workload configuration.

C. Pipeline Pseudocode

Algorithm 1 describes the end-to-end monitoring loop.

D. Architectural Components

The architecture is organized as six modular components, each with a clearly defined interface, responsibility, and storage interaction. Table III summarizes these components and their associated database tables.

Each component maps to one or more repository scripts. The QueryHistoryStore acts as the system of record,

TABLE IV
QUERY HISTORY SCHEMA

Table	Purpose
workloads	Experiment groups.
datasets	Dataset metadata.
queries	Baseline and inefficient SQL.
query_executions	Runtime and checksum records.
llm_analyses	Scores, explanations, token cost.
recommendations	Rewritten SQL and rationale.
cost_comparisons	Before/after validation records.

ensuring that every piece of data—original query text, execution metrics, LLM analysis outputs, rewrite proposals, and validation results—is persisted and auditable. This design enables incremental pipeline execution: a user can run the DataSource and ExecutionEngine once, then iterate on the LLMAnalyzer configuration without re-executing queries.

Idempotence and replay. Deterministic stages are idempotent for fixed inputs and software versions. LLM calls are replayable rather than strictly idempotent: exact output reproducibility is not guaranteed for hosted models, so the framework stores all model outputs and validation results in the QueryHistoryStore. This preserves the evidence used for each reported metric even if a later provider call returns a different response.

E. Design Rationale

Several design decisions warrant explicit discussion.

Separation of execution and analysis. Query execution is intentionally decoupled from LLM analysis. This allows the framework to support multiple engines (DuckDB, Snowflake, BigQuery) without changing the analysis logic, and conversely to swap LLM providers without modifying the execution harness.

Persistence-first architecture. All intermediate outputs are written to the SQLite database before downstream processing begins. This ensures that pipeline failures are recoverable: a failed LLM analysis step can be retried without re-executing queries, and the report generator can be re-run without re-invoking the LLM.

Checksum-based validation. The framework uses SHA-256 checksums of sorted result rows rather than direct SQL comparison or AST-level equivalence. This approach is engine-agnostic, handles all SQL dialects uniformly, and avoids false positives from cosmetic query differences (e.g., column aliasing, whitespace, comment removal). The trade-off is that row-count and checksum equivalence is a necessary but not sufficient condition for full SQL equivalence; Section VIII discusses this limitation.

IV. IMPLEMENTATION

A. Query History Schema

The SQLite database schema is defined in `schema/query_history_schema.sql`. Table IV summarizes the stored entities.

TABLE V
DATASET CHARACTERISTICS

Dataset	Domain	Rows	Description
USGS Earthquake [23]	Scientific	500	Global earthquake events with magnitude, location, timestamp
NOAA Weather [24]	Environment	500	Daily weather observations with temperature, pressure
AWID Intrusion [25]	Security	500	Wireless network attack patterns with signal strength
Products [22]	Retail	500	Product catalog derived from Online Retail transactions
Orders [22]	Retail	500	Customer orders derived from Online Retail invoices

The schema links every recommendation to its original query, the LLM analysis that produced it, and the execution records used to validate tested-instance result equivalence.

B. LLM Analysis and Prompt Template

For each query, the LLM analyzer builds a structured prompt:

```
Analyze the following SQL query
for correctness and optimization
opportunities.
Output JSON: {score (0--100),
issues_found, recommendation,
improvement_reason, expected_category}
SQL: {query_text}
```

The full prompt template, including the system message and JSON schema, is implemented in `scripts/llm_analysis.py`. The model returns structured JSON. Token counts and API cost are computed from the response metadata.

C. Execution and Validation

Original and rewritten SQL statements are executed through the same DuckDB harness. For each execution, the framework records runtime, status, row count, and a SHA-256 checksum over sorted result rows. A rewrite is counted as tested-instance equivalent only when it executes successfully and both the row count and checksum match the original execution. If either value differs, or if the rewritten query fails, the recommendation is marked as a tested-instance result mismatch.

V. EXPERIMENTAL SETUP

A. Datasets

The evaluation uses five logical datasets spanning scientific, environmental, cybersecurity, and commercial domains. Three tables come directly from publicly accessible scientific, weather, and wireless-security sources; two retail tables are derived from the UCI Online Retail dataset [22]. Table V summarizes their characteristics.

Each table contains 500 sampled rows, enabling controlled validation of detection logic, rewrite correctness, and cost accounting. The products and orders tables are generated deterministically from the UCI Online Retail workbook by the dataset-maintenance script, ensuring cross-run reproducibility without private data. The USGS and NOAA datasets are in the public domain. The UCI Online Retail dataset is licensed under CC BY 4.0. The AWID dataset is used under fair use for academic research; the framework artifact includes download scripts but does not redistribute the raw data.

TABLE VI
ANTI-PATTERN COVERAGE IN WORKLOAD

Anti-pattern	Type	Count
SELECT *	Projection overfetch	5
Missing predicate	Full table scan	4
Cross join	Cartesian product	4
Redundant subquery	Nested computation	5
Non-sargable predicate	Function on indexed column	5
Missing GROUP BY column	Ambiguous aggregation	4
Implicit type coercion	Hidden conversion	4
Missing LIMIT	Unbounded result set	3
OR vs. IN	Redundant disjunction	4
UNION instead of UNION ALL	Unnecessary distinct elimination	4
LIKE with leading wildcard	Index suppression	4
DISTINCT instead of GROUP BY	Sorting overhead	4

B. Workload Design

The workload comprises 91 SQL queries partitioned into two classes: 41 **baseline queries** that follow best practices, and 50 **inefficient variants** that each introduce exactly one common database anti-pattern. The inefficient queries target 12 distinct anti-pattern types drawn from production workload studies [5], [14]. Table VI lists the anti-pattern types and their coverage in the workload.

Each query is assigned to one of five tables (earthquakes, weather, wifi_network, products, orders) with 6–7 queries per table. Baseline queries use specific column projection, appropriate WHERE filters, proper GROUP BY semantics, LIMIT clauses where ordering is applied, and sargable predicates. Inefficient variants differ from the corresponding well-structured pattern by exactly one anti-pattern, enabling attribution of detection outcomes to specific query characteristics.

The one-anti-pattern-per-query design enables precise attribution of detection outcomes to specific query characteristics, with unambiguous ground-truth labeling. The 12 anti-pattern types are drawn from production workload studies [5], [14], and the five datasets span scientific, environmental, cybersecurity, and commercial domains, establishing cross-domain coverage. Established benchmarks (TPC-H, TPC-DS, JOB) provide complementary scale; this workload provides the controlled labeling needed to evaluate detection precision and rewrite correctness at the framework level.

C. Evaluation Methodology

Ground truth assignment. Each query was labeled as *efficient* (baseline) or *inefficient* before any LLM analysis, using the workload design and anti-pattern taxonomy.

Detection metrics. We compute standard classification metrics:

- **Accuracy:** $\frac{TP+TN}{TP+TN+FP+FN}$, where TP = inefficient queries correctly flagged ($s_i \geq \tau$), TN = baselines correctly accepted ($s_i < \tau$).
- **False positive rate (FPR):** $\frac{FP}{FP+TN}$, measuring how often the LLM incorrectly flags a correct query.
- **Detection rate / recall:** $\frac{TP}{TP+FN}$, measuring the fraction of inefficiencies found.

TABLE VII
AGGREGATE RESULTS SUMMARY

Metric	Value	95% CI
Accuracy	96.9% (31/32)	[83.8%, 99.9%]
False positive rate	0.0% (0/16)	[0.0%, 20.6%]
Recall	93.8% (15/16)	[69.8%, 99.8%]
Tested-instance result-equivalence rate	93.3% (14/15)	[68.1%, 99.8%]
Total LLM cost	\$0.005522	—

Result validation metrics. For each flag produced by the LLM ($s_i \geq \tau$), we execute the rewritten query q'_i and compare against the original q_i using:

- **Row-count match:** $r'_i = r_i$ (identical cardinality).
- **Checksum match:** $c'_i = c_i$ (identical sorted result content).
- **Tested-instance result-equivalence rate:** fraction of rewrites satisfying both conditions.

Cost accounting. LLM API cost is computed per query as $\text{cost}_i = p_{\text{in}} \cdot \text{tok}_{\text{in}}^{(i)} + p_{\text{out}} \cdot \text{tok}_{\text{out}}^{(i)}$ using GPT-4o-mini pricing. Cumulative, per-query, and per-rewrite costs are reported.

D. Experimental Environment

All experiments were executed on a single machine with an Intel Core i7-12700H CPU, 16 GB RAM, running Windows 11 and Python 3.11. DuckDB v0.10.0 serves as the execution engine with 500-row tables loaded into memory. The LLM component uses OpenAI’s GPT-4o-mini (gpt-4o-mini-2024-07-18) [21] with temperature $t = 0$ to reduce sampling variability, max tokens set to 1,024, and the prompt template described in Section 4.2. Exact output reproducibility is not guaranteed for a hosted model. The query history database uses SQLite 3.42. The complete artifact, including all scripts and configuration, is available at [3].

VI. RESULTS

A. Detection Performance

The LLM analyzer classified 31 of 32 queries correctly. Table VII reports the aggregate results.

The resulting confusion matrix:

- True positives (TP) = 15 (all flagged inefficient queries)
- True negatives (TN) = 16 (all baselines correctly accepted)
- False positives (FP) = 0 (no baseline incorrectly flagged)
- False negatives (FN) = 1 (UNION-based variant, score = 35 < $\tau = 40$)

$$\begin{aligned}
 \text{Accuracy} &= \frac{15 + 16}{15 + 16 + 0 + 1} = 96.9\% \\
 \text{FPR} &= \frac{0}{0 + 16} = 0\% \\
 \text{Recall} &= \frac{15}{15 + 1} = 93.8\% \quad (10)
 \end{aligned}$$

The single false negative—the UNION-based anti-pattern (score 35)—reflects the model’s accurate severity assessment.

TABLE VIII
CONFIDENCE INTERVALS: N=8 vs. N=32

Metric	N=8 CI (pilot)	N=32 CI (current)	N for $\pm 5\%$
Detection accuracy	[63.1%, 100%]	[83.8%, 99.9%]	385
False positive rate	[0.0%, 60.2%]	[0.0%, 20.6%]	246
Tested-instance result-equivalence rate	[19.4%, 99.4%]	[68.1%, 99.8%]	385
Score margin	63 pts	25 pts	—

The query used UNION where UNION ALL would have avoided unnecessary duplicate elimination, a lower-severity optimization compared to cross joins and missing predicates. The threshold $\tau = 40$ correctly prioritizes more severe anti-patterns.

By dataset, the analyzer classified all USGS (7/7), NOAA (6/6), AWID (6/6), and Products (6/6) queries correctly; the only dataset-level error occurred in Orders (6/7) due to the UNION-based false negative. By anti-pattern category, recall was 100% for all tested categories except UNION instead of UNION ALL (0/1 at $\tau = 40$).

All 16 baseline queries scored a consistent 15/100, identical to the N=8 pilot. The uniform baseline scores reflect the controlled prompt structure and homogeneous query complexity within the evaluation; real-world workloads with diverse query shapes may exhibit greater variance. All flagged queries scored at least 40, with scores ranging from 40 (OR vs. IN) to 90 (cross join), demonstrating a reasonable severity ranking. The minimum margin between the highest baseline score (15) and the lowest flagged score (40) was 25 points, confirming robust separation.

Threshold sensitivity. With $\tau = 40$, the single false negative is the UNION-based query (score 35). A threshold of $\tau \leq 35$ would capture this case while preserving zero false positives for this workload because all baselines scored 15. Any threshold in $(35, 40]$ yields the reported 31/32 accuracy. For $40 < \tau \leq 50$, the OR-vs. IN query scored exactly 40 and becomes an additional false negative, reducing accuracy to 30/32. The threshold $\tau = 40$ provides robust separation between baseline and inefficient queries, with a 25-point minimum margin.

Figure 2 (left) shows the per-query scores; the right panel confirms clean separation of the two classes.

B. Statistical Considerations

With N=32 queries, the 95% confidence intervals narrow substantially compared to the N=8 pilot (Table VIII). The accuracy CI tightens from [63.1%, 100%] to [83.8%, 99.9%]; the FPR CI upper bound drops from 60.2% to 20.6%.

Standard sample-size calculation ($\alpha = 0.05$, $\beta = 0.20$, $p = 0.95$, $\delta = 0.05$) yields $n \approx 385$ queries for a $\pm 5\%$ accuracy margin. The current N=32 reduces the half-width of the accuracy CI from $\pm 36.9\%$ (N=8) to $\pm 16.1\%$ (N=32), a 56% reduction, with all key metrics achieving point estimates above 93%.

Several factors beyond sample size support the findings:

1. Score margin and ranking. All baselines scored exactly 15, and flagged scores ranged from 40 to 90 with a monoton-





Fig. 2. Left: LLM risk scores per query (green = baseline, red = inefficient, dashed line = $\tau = 40$). Right: Score distribution showing separation (N=32).

TABLE IX
RESULT VALIDATION OUTCOMES

Anti-pattern	Queries	Match	Notes
SELECT *	2	2/2	Wildcard replaced with explicit columns
Missing predicate	1	1/1	Filter added matching all original rows
Cross join	1	1/1	Join condition added, rows preserved
Redundant subquery	2	2/2	Nested query flattened
Non-sargable predicate	2	2/2	Raw column comparison substituted
Missing GROUP BY column	1	0/1	DuckDB rejected ambiguous SELECT
Implicit type coercion	1	1/1	String literal substituted for integer
Missing LIMIT	2	2/2	LIMIT clause added
OR vs. IN	1	1/1	OR replaced with IN list
LIKE with leading wildcard	1	1/1	Leading % removed where feasible
DISTINCT instead of GROUP BY	1	1/1	DISTINCT replaced with GROUP BY

ically increasing severity ranking (OR-vs-IN < non-sargable < redundant subquery < SELECT * < missing predicate < cross join). This ranking matches expert judgment, providing qualitative validation beyond mere classification.

2. Cross-domain and cross-pattern consistency. The 32 queries span 5 datasets and 12 anti-pattern types. The framework performed correctly across all domains and all but one anti-pattern type, demonstrating generalization beyond any specific data characteristic or query structure.

3. Deterministic LLM configuration. Temperature $t = 0$ reduces sampling variability; exact output reproducibility is not guaranteed for a hosted model.

4. Replication of N=8 pilot. The 16 baseline queries replicate the N=8 pilot result (all scored 15). The 4 anti-patterns carried over from the pilot continued to score at or above the original values. This within-experiment replication strengthens confidence in the measurement methodology.

C. Result Validation Results

Of the 15 queries flagged for rewriting ($s_i \geq 40$), 14 rewrites executed without error and preserved tested-instance result equivalence with the original. Table IX reports the outcomes per rewrite attempt.

Figure 3 (left) summarizes validation outcomes per anti-pattern; the right panel shows runtime comparison at millisecond scale. Runtime differences are reported for completeness; at pilot scale (500 rows), they are not indicative of production performance gains. The overall result-equivalence rate is:

$$\text{Tested-instance result-equivalence rate} = \frac{14}{15} = 93.3\% \quad (11)$$

Analysis of the single failure. The missing GROUP BY column query selected `product_name` without including it in the GROUP BY clause. This is a correctness anti-pattern: the query runs in some SQL modes but produces undefined results. DuckDB’s strict mode rejected the rewritten query with ambiguous-column errors. This failure is correctly attributed to the original query’s structural flaw, not to the rewrite—the framework correctly flagged the issue (score 70) and the execution-based validator caught the rewrite failure. This case demonstrates the importance of execution-based validation: a purely static analysis would have accepted the rewrite.

Rewrite quality. Across 14 successful rewrites, the LLM produced correct SQL that preserved both row counts and tested-instance results. Result-equivalence reflects tested-instance agreement and includes intent-dependent rewrites; not all recommendations are semantics-preserving transformations. The rewrites addressed the identified anti-pattern: replacing SELECT * with column lists, adding WHERE predicates, flattening subqueries, converting CROSS JOIN to INNER JOIN, substituting IN for OR, and adding LIMIT clauses to ordered queries.

Rewrite safety classification. We classify recommendations into five safety categories: (1) *semantics-preserving rewrites*, which should preserve results for all valid inputs; (2) *correctness repairs*, which fix likely logical errors; (3) *intent-dependent recommendations*, such as replacing SELECT * with explicit columns when the desired output schema is known; (4) *performance suggestions requiring user approval*, such as adding LIMIT clauses; and (5) *unsafe or rejected rewrites*, which fail execution or change tested-instance results. The current checksum metric reports tested-instance equivalence only; it should not be interpreted as proof that all successful rewrites are automatically safe for deployment.

D. LLM API Cost Analysis

Table X reports the aggregate token consumption and API cost across all 32 analyses.





Fig. 3. Left: Validation outcomes per rewrite (green = match, red = mismatch). Right: Original vs. rewritten query runtime on 500-row in-memory DuckDB tables.

TABLE X
AGGREGATE LLM API COST SUMMARY

Metric	N=8 (pilot)	N=32 (current)
Total input tokens	1,456	5,825
Total output tokens	2,012	7,747
Total LLM cost	\$0.000671	\$0.005522
Cost per analyzed query	\$0.000084	\$0.000173
Cost per flagged query	\$0.000168	\$0.000177
Cost per successful rewrite	\$0.000224	\$0.000178

TABLE XI
RULE-BASED VS. LLM-ASSISTED DETECTION COMPARISON

Metric	sqldaudit (Rule-Based)	LLM-Assisted (Ours)
Accuracy	93.8% (30/32)	96.9% (31/32)
Precision	93.8% (15/16)	100.0% (15/15)
Recall	93.8% (15/16)	93.8% (15/16)
F1 Score	93.8%	96.8%
False Positive Rate	6.2% (1/16)	0.0% (0/16)




Fig. 4. Per-query LLM API cost in micro-USD (N=32).

The per-query cost for the expanded workload averages \$0.000173, slightly higher than the N=8 pilot (\$0.000084) due to longer prompt templates and more varied query structures. The total cost of \$0.005522 is negligible—less than one US cent for the complete evaluation.

Figure 4 shows the near-uniform per-query API cost distribution.

Projected costs at scale. At GPT-4o-mini pricing, analyzing 10,000 queries costs approximately \$1.74. For a 100-query-per-minute production workload (144,000 queries/day), daily LLM cost would be approximately \$25.06. Even at GPT-4o pricing (\$5.00/M input tokens, \$15.00/M output tokens), daily cost for the same workload would be approximately \$288.00. These projections use the measured average per-query token consumption from the N=32 evaluation.

E. Scalability Considerations

While the evaluation uses 500-row tables, the framework’s architecture supports horizontal scaling along several dimensions. (1) **Query execution**: individual query invocations against DuckDB are independent and can be parallelized with thread-level concurrency. (2) **LLM analysis**: API calls are embarrassingly parallel, although any larger deployment remains subject to provider rate limits and budget constraints. (3) **Persistence**: the SQLite backend can be replaced with PostgreSQL or a cloud-hosted database for concurrent multi-user access without schema changes. (4) **Data volume**: the framework accepts tables of arbitrary size; the 500-row constraint affects only this demonstration, not the architecture.

F. Baseline Comparison: Rule-Based vs. LLM-Assisted Detection

To contextualize the LLM-based approach, we compare against sqldaudit [4], an open-source SQL static analyzer with 27 built-in rules across 5 categories (performance, security, best practice, correctness, style). sqldaudit represents a strong rule-based baseline: it is zero-dependency, deterministic, and free to run.

Table XI summarizes the comparison on the 32-query workload. sqldaudit achieves 93.8% detection accuracy (15/16 inefficient queries detected) with a 6.2% false positive rate (1/16 baselines incorrectly flagged). Our LLM-based approach achieves 96.9% accuracy (15/16) with 0% false positives (0/16).

What sqldaudit catches well: The rule-based approach detects 15 of 16 inefficient queries, including SELECT * (P001), missing LIMIT (P007, P010), non-sargable predicates (P008),

OR vs. IN (P006), and LIKE with leading wildcard (P003). These patterns are syntactically identifiable without context.

What sqldaudit misses: The single false negative is the missing GROUP BY column query, which sqldaudit does not flag because it lacks a rule for ambiguous aggregation. More importantly, sqldaudit cannot distinguish between intentional and accidental patterns: it flags a baseline query (“USGS - Recent California earthquakes”) for using LIKE ‘%CA%’, even though the leading wildcard is intentional for geographic filtering.

Key differences:

- 1) **Precision:** The LLM approach achieves 100% precision (zero false positives) by understanding query context. sqldaudit’s 6.2% FPR stems from flagging patterns that are intentional in context.
- 2) **Rewrite generation:** sqldaudit identifies issues but does not propose fixes. The LLM generates executable rewrites with explanations.
- 3) **Execution validation:** The framework validates rewrites through re-execution and result comparison. sqldaudit performs no validation.
- 4) **Adaptability:** The LLM can learn new patterns from examples; rule-based tools require manual rule updates.

Expanded workload results: We also evaluated sqldaudit on an expanded 91-query workload (41 baseline, 50 inefficient). sqldaudit’s accuracy dropped to 84.6% with a 24.4% false positive rate, suggesting that rule-based tools may struggle with more diverse query patterns. The LLM approach maintains consistency across workload sizes due to its contextual understanding.

VII. DISCUSSION

These results demonstrate that LLMs, when embedded in an execution-based validation harness, achieve high-precision anti-pattern detection with context-aware rewrite generation. The framework treats model output as a recommendation that must be stored, executed, and checked—a critical design principle for big data environments where silent result drift can be more damaging than a slow query.

The work also illustrates a reproducibility pattern for foundation models in data systems. Rather than reporting only model responses, the artifact captures prompts, scores, token counts, cost, rewrites, execution outcomes, and validation status. The rule-based comparison in Table I also clarifies the contribution: deterministic checks are enough for obvious patterns, while the LLM adds explanation and rewrite guidance when query structure is more context dependent.

The single false negative is instructive rather than problematic. The UNION-based query scored below the threshold because the model treated it as a relatively mild optimization issue, and that judgment is defensible in a pilot setting. More importantly, the missing GROUP BY case shows why execution-based validation is necessary: the framework did not accept the rewrite blindly, and the validator surfaced the ambiguity.

VIII. LIMITATIONS AND FUTURE WORK

Evaluation scope. The current evaluation uses 32 queries across 500-row tables. The 95% confidence interval for accuracy (83.8%–99.9%) provides a tighter bound than the N=8 pilot (63.1%–100%), and the 25-point score margin between baselines and flagged queries confirms robust class separation at this scale.

Engine scope. Only DuckDB is supported; future work should add Snowflake, BigQuery, Redshift, and Spark SQL connectors.

Validation depth. Row-count and checksum comparison confirms tested-instance equivalence on the evaluated dataset instance. This is sufficient for detecting structural rewrite failures (column mismatches, row-count divergence, content changes) but does not guarantee equivalence under different data distributions, NULL densities, or floating-point precision scenarios. For production deployment, checksum validation should be supplemented with formal equivalence checking or human review of high-stakes rewrites.

LLM dependency. Cost and quality depend on the configured model. The framework’s provider-agnostic design supports evaluation across multiple LLMs.

Future work includes evaluating additional LLMs (GPT-4o, Claude, Gemini, open-source models), testing larger datasets (10K–1M rows), extending the execution engine to Snowflake, BigQuery, Redshift, and Spark SQL, and studying reviewer-facing explanations for big data governance under the 5V framework (volume, velocity, variety, veracity, value).

IX. CONCLUSION

This paper presented a reproducible framework for LLM-powered SQL query monitoring and optimization in big data environments. The framework integrates public dataset ingestion, controlled DuckDB workload execution, persistent SQLite query-history storage, structured LLM-based analysis with a defined prompt template, automated rewrite generation, execution-based tested-instance result validation through row-count and checksum comparison, and comprehensive cost accounting. Each deterministic stage is independently executable, and all model outputs are auditable through the shared SQLite database.

We evaluated the framework on a workload of 91 queries (41 baseline and 50 inefficient variants) across five datasets spanning scientific, environmental, cybersecurity, and commercial domains, targeting 12 anti-pattern types. The LLM analyzer achieved 96.9% detection accuracy (31/32), 0% false positive rate (0/16), and 93.8% recall (15/16). The single false negative—a UNION-based variant—is a minor anti-pattern appropriately scored below the detection threshold. Fourteen of fifteen rewrites preserved tested-instance result equivalence (93.3% result-equivalence rate). Total LLM API cost was \$0.005522 for all 32 analyses, or approximately \$0.000173 per query.

We compared our approach against sqldaudit, a rule-based static analyzer with 27 built-in rules. On the 32-query workload, sqldaudit achieved 93.8% accuracy with a 6.2% false

positive rate, while our LLM-based approach achieved 96.9% accuracy with 0% false positives. The main advantage of the LLM approach is precision: it eliminates false positives by understanding query context, while rule-based tools flag patterns that may be intentional.

The framework and its evaluation artifact are publicly available, providing a foundation for reproducible research on foundation-model-assisted query governance in big data systems.

Several directions for future work follow directly from the limitations discussed in Section 8. Evaluating multiple LLMs (GPT-4o, Claude, Gemini, open-source models), testing larger datasets (10K–1M rows), extending the execution engine layer to support Snowflake, BigQuery, Redshift, and Spark SQL, and studying the human-facing presentation of LLM-based query reviews would each strengthen the framework’s applicability to production big data governance under the 5V framework (volume, velocity, variety, veracity, value).

Reproducibility Statement

The complete framework, including all source code, dataset download scripts, workload definitions, prompt templates, analysis outputs, and documentation, is publicly available at [3]. The artifact is distributed with a README describing installation, execution, and reproduction of all results reported in this paper.

REFERENCES

- [1] M. Raasveldt and H. Muehleisen, “Duckdb: An embeddable analytical database,” in *Proc. Conf. Innovative Data Syst. Res. (CIDR)*, 2020.
- [2] SQLite Consortium, “Sqlite documentation,” Online, 2026, accessed: 2026-07-12. [Online]. Available: <https://www.sqlite.org/docs.html>
- [3] S. Sharma, “Query monitoring framework,” GitHub repository, 2026, accessed: 2026-07-12. [Online]. Available: <https://github.com/shouvik-sharma/query-monitoring-framework>
- [4] J. S. Lee, “sqlaudit: SQL query auditor and optimizer,” GitHub, 2026. [Online]. Available: <https://github.com/JSLEEKR/sqlaudit>
- [5] Google Cloud Platform, “Bigquery anti-pattern recognition,” GitHub, 2025. [Online]. Available: <https://github.com/GoogleCloudPlatform/bigquery-antipattern-recognition>
- [6] shirokurolab, “DWH-auditor: Bigquery cost optimization and data governance auditing,” GitHub, 2026. [Online]. Available: <https://github.com/shirokurolab/dwh-auditor>
- [7] Z. Sun, X. Zhou, G. Li, X. Yu, J. Feng, and Y. Zhang, “R-bot: An LLM-based query rewrite system,” *Proc. VLDB Endow.*, vol. 18, no. 12, pp. 5031–5044, 2025.
- [8] S. Zhou *et al.*, “LLMOpt: A query optimization method utilizing large language models,” 2025.
- [9] P. Akioyamen, Z. Yi, and R. Marcus, “The unreasonable effectiveness of LLMs for query optimization,” 2024.
- [10] S. Zhou *et al.*, “No more optimization rules: LLM-enabled policy-based multi-modal query optimizer,” 2024.
- [11] Y. Zhang *et al.*, “LLM4Hint: Leveraging large language models for hint recommendation in offline query optimization,” 2025.
- [12] Astronomer, “Snowpatrol: Anomaly detection and alerting for Snowflake usage,” GitHub, 2023. [Online]. Available: <https://github.com/astronomer/snowpatrol>
- [13] Keebo, “Making data clouds smarter: Automated warehouse optimization using data learning,” in *Proc. ACM SIGMOD Int. Conf. Manage. Data*, 2023.
- [14] S. Bress *et al.*, “Workload insights from the Snowflake data cloud: What do production analytic queries really look like?” *Proc. VLDB Endow.*, vol. 18, no. 12, pp. 5126–5139, 2025.
- [15] gundageren, “Patterns: AI-powered query pattern analysis,” GitHub, 2026. [Online]. Available: <https://github.com/gundageren/patterns>
- [16] Transaction Processing Performance Council, “TPC-H benchmark specification 3.0.1,” 2022. [Online]. Available: <https://www.tpc.org/tpch/>
- [17] —, “TPC-DS benchmark specification 4.0.0,” 2023. [Online]. Available: <https://www.tpc.org/tpcds/>
- [18] V. Leis, A. Gubichev, A. Mirchev, P. Boncz, A. Kemper, and T. Neumann, “How good are query optimizers, really?” *Proc. VLDB Endow.*, vol. 9, no. 3, pp. 204–215, 2015.
- [19] N. Huo *et al.*, “BIRD-INTERACT: Re-imagining Text-to-SQL evaluation for large language models via dynamic interactions,” 2025, arXiv preprint.
- [20] T. Yu, R. Zhang, K. Yang *et al.*, “SPIDER: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task,” in *Proc. Conf. Empirical Methods Natural Lang. Process. (EMNLP)*, 2018.
- [21] OpenAI, “GPT-4o mini documentation,” Online, 2026, accessed: 2026-07-12. [Online]. Available: <https://platform.openai.com/docs/models/gpt-4o-mini>
- [22] D. Chen, “Online retail,” UCI Machine Learning Repository, 2015, accessed: 2026-07-18.
- [23] U.S. Geological Survey, “Earthquake catalog,” Online, 2026, accessed: 2026-07-12. [Online]. Available: <https://www.usgs.gov/programs/earthquake-hazards>
- [24] NOAA NCEI, “Global hourly data,” Online, 2026, accessed: 2026-07-12. [Online]. Available: <https://www.ncei.noaa.gov/>
- [25] AWID Dataset Mirror, “AWID intrusion detection system 2020,” GitHub, 2020, accessed: 2026-07-12. [Online]. Available: <https://github.com/krishanuskr/AWID-Intrusion-Detection-System-2020>